

openmic

December 17, 2019

```
[1]: import librosa as lb
import librosa.display
import pandas as pd
import scipy
import json
import numpy as np
import sklearn
from sklearn.metrics import classification_report
from sklearn.model_selection import train_test_split
import os
from pylab import plot, show, figure, imshow, xlim, ylim, title
import matplotlib.pyplot as plt
import keras
from keras.utils import np_utils
from keras import layers
from keras import models
```

Using TensorFlow backend.

```
[2]: #CONSTANTS

DATA_DIR = "openmic-2018/"
CATEGORY_COUNT = 8

[3]: df = pd.read_csv('openmic-2018/openmic-2018-aggregated-labels.csv')
del df['relevance']
del df['num_responses']

[4]: labels = df.values
labels

[4]: array([[ '000046_3840', 'clarinet'],
          [ '000046_3840', 'flute'],
          [ '000046_3840', 'trumpet'],
          ...,
          [ '155311_453120', 'saxophone'],
          [ '155311_453120', 'trumpet'],
          [ '155311_453120', 'trombone']], dtype=object)
```

```
[5]: y, sr = lb.load(DATA_DIR + 'audio/000/000135_483840.ogg')
S = lb.feature.melspectrogram(y=y, sr=sr)


S_dB = lb.power_to_db(S, ref=0) # 10 * log10(S / ref)

print(y.shape)
print(sr)
print(S.shape)
print(S_dB.shape)
```

```
(220544,)
22050
(128, 431)
(128, 431)
```

```
[6]: librosa.display.specshow(S_dB, x_axis='s', y_axis='mel')
plt.colorbar(format='%+2.0f dB')
```

```
[6]: <matplotlib.colorbar.Colorbar at 0x234e4671278>
```



```
[7]: OPENMIC = np.load(os.path.join(DATA_DIR, 'openmic-2018.npz'), allow_pickle=True)
print(list(OPENMIC.keys()))
```

```
['X', 'Y_true', 'Y_mask', 'sample_key']
```

```
[8]: X, Y_true, Y_mask, sample_key = OPENMIC['X'], OPENMIC['Y_true'],  
      ↪OPENMIC['Y_mask'], OPENMIC['sample_key']  
      #print(X.shape)  
      #X = []  
      #print(len(sample_key))  
      #for key in sample_key:  
      #    key_dir = key[:3]  
      #    y, sr = lb.load(DATA_DIR + 'audio/' + key_dir + '/' + key + '.ogg')  
      #    X.append(lb.feature.melspectrogram(y=y, sr=sr))  
      #    print(len(X))
```

```
[9]: with open(os.path.join(DATA_DIR, 'class-map.json'), 'r') as f:  
      class_map = json.load(f)
```

```
[21]: split_train, split_test, X_train, X_test, Y_true_train, Y_true_test,  
      ↪Y_mask_train, Y_mask_test = train_test_split(sample_key, X, Y_true, Y_mask)  
      split_val, split_test, X_val, X_test, Y_true_val, Y_true_test, Y_mask_val,  
      ↪Y_mask_test = train_test_split(split_test, X_test, Y_true_test, Y_mask_test,  
      ↪test_size=0.5)  
  
      train_set = np.asarray(set(split_train))  
      test_set = np.asarray(set(split_test))  
      print('# Train: {}, # Val: {}, # Test: {}'.format(len(split_train),  
      ↪len(split_test), len(split_val)))
```

```
# Train: 15000, # Val: 2500, # Test: 2500
```

```
[23]: print(X_train.shape)  
      print(X_val.shape)  
      print(X_test.shape)
```

```
(15000, 10, 128)
```

```
(2500, 10, 128)
```

```
(2500, 10, 128)
```

```
[27]: THRESHOLD = 0.3  
  
      # This dictionary will include the classifiers for each model  
      mymodels = dict()  
  
      # We'll iterate over all instrument classes, and fit a model for each one  
      # After training, we'll print a classification report for each instrument  
      for instrument in class_map:  
  
          # Map the instrument name to its column number  
          inst_num = class_map[instrument]
```

```

# Step 1: sub-sample the data

# First, we need to select down to the data for which we have annotations
# This is what the mask arrays are for
train_inst = Y_mask_train[:, inst_num]
val_inst = Y_mask_val[:, inst_num]
test_inst = Y_mask_test[:, inst_num]

# Here, we're using the Y_mask_train array to slice out only the training
→examples
# for which we have annotations for the given class
X_train_inst = X_train[train_inst]
X_val_inst = X_val[val_inst]

# Step 3: simplify the data by averaging over time

# Let's arrange the data for a sklearn Random Forest model
# Instead of having time-varying features, we'll summarize each track by
→its mean feature vector over time
X_train_inst_sklearn = np.mean(X_train_inst, axis=1)

# Again, we slice the labels to the annotated examples
# We threshold the label likelihoods at 0.5 to get binary labels
Y_true_train_inst = Y_true_train[train_inst, inst_num] >= THRESHOLD
Y_true_val_inst = Y_true_val[val_inst, inst_num] >= THRESHOLD

# Repeat the above slicing and dicing but for the test set
X_test_inst = X_test[test_inst]
X_test_inst_sklearn = np.mean(X_test_inst, axis=1)
Y_true_test_inst = Y_true_test[test_inst, inst_num] >= THRESHOLD

X_train_inst = X_train_inst.astype('float32')
X_val_inst = X_val_inst.astype('float32')
X_train_inst_sklearn = X_train_inst_sklearn.astype('float32')
X_train_inst_sklearn = lb.util.normalize(X_train_inst_sklearn)
# X_train_inst = S_dB
print(X_train_inst.shape)
shape = X_train_inst.shape
X_train_inst = X_train_inst.reshape(shape[0],1, shape[1], shape[2])
shape = X_val_inst.shape
X_val_inst = X_val_inst.reshape(shape[0],1, shape[1], shape[2])
shape = X_test_inst.shape
X_test_inst = X_test_inst.reshape(shape[0],1, shape[1], shape[2])
#X_train_inst = X_train_inst.reshape(1,1,431,128)
print(X_train_inst.shape)

```

```

print(Y_true_train_inst[0])
# Step 3.
# Initialize a new classifier
import keras, os
from keras.models import Sequential
from keras.layers import Dense, Conv2D, MaxPool2D, Flatten
from keras.preprocessing.image import ImageDataGenerator
import numpy as np
model = models.Sequential()

# model.add(layers.Conv2D(filters=8, kernel_size=(3,3), activation='relu',
→input_shape=(10,128,1,)))
model.
→add(Conv2D(input_shape=(1,10,128), data_format="channels_first", filters=64, kernel_size=(3,3)
→activation="relu"))
model.add(Conv2D(filters=32, kernel_size=(3,3), padding="same",
→activation="relu"))
model.add(MaxPool2D(pool_size=(3,3), strides=(2,2)))
model.add(Conv2D(filters=128, kernel_size=(3,3), padding="same",
→activation="relu"))
model.add(Conv2D(filters=128, kernel_size=(3,3), padding="same",
→activation="relu"))
model.add(MaxPool2D(pool_size=(2,2), strides=(2,2)))
model.add(Conv2D(filters=256, kernel_size=(3,3), padding="same",
→activation="relu"))
model.add(Conv2D(filters=256, kernel_size=(3,3), padding="same",
→activation="relu"))
model.add(Conv2D(filters=256, kernel_size=(3,3), padding="same",
→activation="relu"))
model.add(MaxPool2D(pool_size=(2,2), strides=(2,2)))
model.add(Conv2D(filters=512, kernel_size=(3,3), padding="same",
→activation="relu"))
model.add(Conv2D(filters=512, kernel_size=(3,3), padding="same",
→activation="relu"))
model.add(Conv2D(filters=512, kernel_size=(3,3), padding="same",
→activation="relu"))
model.add(layers.Flatten())
model.add(layers.Dense(units=4096, activation='relu'))
model.add(layers.Dense(units=4096, activation='relu'))
model.add(layers.Dense(units=1, activation='sigmoid'))

model.compile(loss='binary_crossentropy',
              optimizer=keras.optimizers.Adam(lr=0.00001),
              metrics = ['accuracy'])

# model.summary()

```

```

# Step 4.
history = model.fit(X_train_inst, Y_true_train_inst , epochs=10,
→batch_size=64, validation_data=(X_val_inst, Y_true_val_inst))

plt.figure(figsize=(9,4))

plt.subplot(1,2,1)
plt.plot(history.history['acc'])
plt.plot(history.history['val_acc'])
plt.title('Model accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train accuracy', 'Validation accuracy'], loc='upper left')

plt.subplot(1,2,2)
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train loss', 'Validation loss'], loc='upper left')
plt.show()

loss, acc = model.evaluate(X_test_inst, Y_true_test_inst)
print('Test loss: {}'.format(loss))
print('Test accuracy: {:.2%}'.format(acc))

# Step 5.
# Finally, we'll evaluate the model on both train and test
Y_pred_train = model.predict(X_train_inst)
Y_pred_test = model.predict(X_test_inst)
Y_pred_train_bool = Y_pred_train > THRESHOLD #THRESHOLD (should be lower
→than 0.5)
Y_pred_test_bool = Y_pred_test > THRESHOLD #THRESHOLD (should be lower than
→0.5)

print(Y_pred_train[0])
print('-' * 52)
print(instrument)
print('\tTRAIN')
print(classification_report(Y_true_train_inst, Y_pred_train_bool))

print(Y_true_train_inst[3])
print(Y_pred_train[3])
print('\tTEST')
print(classification_report(Y_true_test_inst, Y_pred_test_bool))
sum = 0
# for i, prob in enumerate(Y_pred_train):

```

```

#     print (i)
#     print (prob)
#     sum += prob
#     print(sum)
# Store the classifier in our dictionary
mymodels[instrument] = model

```

(1566, 10, 128)

(1566, 1, 10, 128)

False

Train on 1566 samples, validate on 249 samples

Epoch 1/10

1566/1566 [=====] - 90s 57ms/step - loss: 0.5710 - acc: 0.7625 - val_loss: 0.5417 - val_acc: 0.7671

Epoch 2/10

1566/1566 [=====] - 87s 56ms/step - loss: 0.5380 - acc: 0.7625 - val_loss: 0.5206 - val_acc: 0.7671

Epoch 3/10

1566/1566 [=====] - 88s 56ms/step - loss: 0.5110 - acc: 0.7625 - val_loss: 0.4943 - val_acc: 0.7671

Epoch 4/10

1566/1566 [=====] - 88s 56ms/step - loss: 0.4703 - acc: 0.7822 - val_loss: 0.4399 - val_acc: 0.7992

Epoch 5/10

1566/1566 [=====] - 87s 55ms/step - loss: 0.4196 - acc: 0.8084 - val_loss: 0.4243 - val_acc: 0.7992

Epoch 6/10

1566/1566 [=====] - 86s 55ms/step - loss: 0.4006 - acc: 0.8199 - val_loss: 0.4090 - val_acc: 0.8112

Epoch 7/10

1566/1566 [=====] - 86s 55ms/step - loss: 0.3760 - acc: 0.8340 - val_loss: 0.4124 - val_acc: 0.8072

Epoch 8/10


1566/1566 [=====] - 86s 55ms/step - loss: 0.3591 - acc: 0.8442 - val_loss: 0.4015 - val_acc: 0.8273

Epoch 9/10

1566/1566 [=====] - 86s 55ms/step - loss: 0.3615 - acc: 0.8436 - val_loss: 0.3998 - val_acc: 0.8072

Epoch 10/10

1566/1566 [=====] - 85s 55ms/step - loss: 0.3551 - acc: 0.8359 - val_loss: 0.4271 - val_acc: 0.7791



256/256 [=====] - 2s 10ms/step
 Test loss: 0.4729207083582878
 Test accuracy: 75.00%
 [0.7572417]

 accordion

TRAIN					
	precision	recall	f1-score	support	
False	0.96	0.71	0.82	1194	
True	0.49	0.92	0.64	372	
accuracy			0.76	1566	
macro avg	0.73	0.81	0.73	1566	
weighted avg	0.85	0.76	0.77	1566	

False

[0.6689949]


TEST					
	precision	recall	f1-score	support	
False	0.90	0.66	0.76	197	
True	0.40	0.75	0.52	59	
accuracy			0.68	256	
macro avg	0.65	0.70	0.64	256	
weighted avg	0.78	0.68	0.70	256	

(1669, 10, 128)


```

(1669, 1, 10, 128)
True
Train on 1669 samples, validate on 279 samples
Epoch 1/10
1669/1669 [=====] - 102s 61ms/step - loss: 0.6316 -
acc: 0.6729 - val_loss: 0.6269 - val_acc: 0.6559
Epoch 2/10
1669/1669 [=====] - 93s 56ms/step - loss: 0.6120 - acc:
0.6729 - val_loss: 0.5936 - val_acc: 0.6559
Epoch 3/10
1669/1669 [=====] - 93s 56ms/step - loss: 0.5465 - acc:
0.6998 - val_loss: 0.5156 - val_acc: 0.7993
Epoch 4/10
1669/1669 [=====] - 93s 56ms/step - loss: 0.4749 - acc:
0.7687 - val_loss: 0.4455 - val_acc: 0.7885
Epoch 5/10
1669/1669 [=====] - 93s 55ms/step - loss: 0.4302 - acc:
0.7801 - val_loss: 0.4325 - val_acc: 0.7885
Epoch 6/10
1669/1669 [=====] - 101s 60ms/step - loss: 0.4065 -
acc: 0.7945 - val_loss: 0.4316 - val_acc: 0.7957
Epoch 7/10
1669/1669 [=====] - 101s 61ms/step - loss: 0.4036 -
acc: 0.8005 - val_loss: 0.4614 - val_acc: 0.7742
Epoch 8/10
1669/1669 [=====] - 107s 64ms/step - loss: 0.4192 -
acc: 0.7927 - val_loss: 0.4812 - val_acc: 0.7240
Epoch 9/10
1669/1669 [=====] - 108s 65ms/step - loss: 0.3849 -
acc: 0.8107 - val_loss: 0.4528 - val_acc: 0.7706
Epoch 10/10
1669/1669 [=====] - 102s 61ms/step - loss: 0.3679 -
acc: 0.8250 - val_loss: 0.4608 - val_acc: 0.7706

```



270/270 [=====] - 3s 10ms/step
 Test loss: 0.4562561829884847
 Test accuracy: 77.78%
 [0.9145529]

 banjo

TRAIN				
	precision	recall	f1-score	support
False	0.90	0.85	0.87	1123
True	0.72	0.80	0.76	546
accuracy			0.83	1669
macro avg	0.81	0.83	0.82	1669
weighted avg	0.84	0.83	0.84	1669

False

[0.00839368]

TEST				
	precision	recall	f1-score	support
False	0.86	0.79	0.82	180
True	0.64	0.73	0.68	90
accuracy			0.77	270
macro avg	0.75	0.76	0.75	270
weighted avg	0.78	0.77	0.78	270

(1401, 10, 128)

(1401, 1, 10, 128)
False
Train on 1401 samples, validate on 258 samples
Epoch 1/10

```
KeyboardInterrupt                                Traceback (most recent call
last)

<ipython-input-27-0129277914b6> in <module>
    91     # model.summary()
    92     # Step 4.
--> 93     history = model.fit(X_train_inst,Y_true_train_inst , epochs=10,
batch_size=64, validation_data=(X_val_inst,Y_true_val_inst))
    94
    95     plt.figure(figsize=(9,4))

~\Documents\Conda\lib\site-packages\keras\engine\training.py in
fit(self, x, y, batch_size, epochs, verbose, callbacks, validation_split,
validation_data, shuffle, class_weight, sample_weight, initial_epoch,
steps_per_epoch, validation_steps, **kwargs)
    1037                                     initial_epoch=initial_epoch,
    1038
-> steps_per_epoch=steps_per_epoch,
-> 1039
validation_steps=validation_steps)
    1040
    1041     def evaluate(self, x=None, y=None,

~\Documents\Conda\lib\site-packages\keras\engine\training_arrays.py in
fit_loop(model, f, ins, out_labels, batch_size, epochs, verbose, callbacks,
val_f, val_ins, shuffle, callback_metrics, initial_epoch, steps_per_epoch,
validation_steps)
    197                                     ins_batch[i] = ins_batch[i].toarray()
    198
--> 199                                     outs = f(ins_batch)
    200                                     outs = to_list(outs)
    201                                     for l, o in zip(out_labels, outs):

~\Documents\Conda\lib\site-packages\keras\backend\tensorflow_backend.py
in __call__(self, inputs)
```

```

2713             return self._legacy_call(inputs)
2714
-> 2715         return self._call(inputs)
2716     else:
2717         if py_any(is_tensor(x) for x in inputs):

~\Documents\Conda\lib\site-packages\keras\backend\tensorflow_backend.py
-> in _call(self, inputs)
    2673         fetched = self._callable_fn(*array_vals,
-> run_metadata=self.run_metadata)
    2674     else:
-> 2675         fetched = self._callable_fn(*array_vals)
    2676         return fetched[:len(self.outputs)]
    2677

~\Documents\Conda\lib\site-packages\tensorflow\python\client\session.py
-> in __call__(self, *args, **kwargs)
    1437         ret = tf_session.TF_SessionRunCallable(
    1438             self._session._session, self._handle, args, status,
-> 1439             run_metadata_ptr)
    1440         if run_metadata:
    1441             proto_data = tf_session.TF_GetBuffer(run_metadata_ptr)

```

KeyboardInterrupt:

```
[ ]: print(X_train_inst_sklearn)
      print(Y_pred_train)
```

```
[ ]:
```