

# openmic

December 17, 2019

```
[1]: import librosa as lb
import librosa.display
import pandas as pd
import scipy
import json
import numpy as np
import sklearn
from sklearn.metrics import classification_report
from sklearn.model_selection import train_test_split
import os
from pylab import plot, show, figure, imshow, xlim, ylim, title
import matplotlib.pyplot as plt
import keras
from keras.utils import np_utils
from keras import layers
from keras import models
```

Using TensorFlow backend.

```
[2]: #CONSTANTS

DATA_DIR = "openmic-2018/"
CATEGORY_COUNT = 8

[3]: df = pd.read_csv('openmic-2018/openmic-2018-aggregated-labels.csv')
del df['relevance']
del df['num_responses']

[4]: labels = df.values
labels

[4]: array([[ '000046_3840', 'clarinet'],
            [ '000046_3840', 'flute'],
            [ '000046_3840', 'trumpet'],
            ...,
            [ '155311_453120', 'saxophone'],
            [ '155311_453120', 'trumpet'],
            [ '155311_453120', 'trombone']], dtype=object)
```

```
[5]: y, sr = lb.load(DATA_DIR + 'audio/000/000135_483840.ogg')
S = lb.feature.melspectrogram(y=y, sr=sr)


S_dB = lb.power_to_db(S, ref=0) # 10 * log10(S / ref)

print(y.shape)
print(sr)
print(S.shape)
print(S_dB.shape)
```

```
(220544,)
22050
(128, 431)
(128, 431)
```

```
[6]: librosa.display.specshow(S_dB, x_axis='s', y_axis='mel')
plt.colorbar(format='%+2.0f dB')
```

```
[6]: <matplotlib.colorbar.Colorbar at 0x234e4671278>
```



```
[7]: OPENMIC = np.load(os.path.join(DATA_DIR, 'openmic-2018.npz'), allow_pickle=True)
print(list(OPENMIC.keys()))
```

```
['X', 'Y_true', 'Y_mask', 'sample_key']
```

```
[8]: X, Y_true, Y_mask, sample_key = OPENMIC['X'], OPENMIC['Y_true'],  
      ↪OPENMIC['Y_mask'], OPENMIC['sample_key']  
      #print(X.shape)  
      #X = []  
      #print(len(sample_key))  
      #for key in sample_key:  
      #    key_dir = key[:3]  
      #    y, sr = lb.load(DATA_DIR + 'audio/' + key_dir + '/' + key + '.ogg')  
      #    X.append(lb.feature.melspectrogram(y=y, sr=sr))  
      #    print(len(X))
```

```
[9]: with open(os.path.join(DATA_DIR, 'class-map.json'), 'r') as f:  
      class_map = json.load(f)
```

```
[21]: split_train, split_test, X_train, X_test, Y_true_train, Y_true_test,  
      ↪Y_mask_train, Y_mask_test = train_test_split(sample_key, X, Y_true, Y_mask)  
      split_val, split_test, X_val, X_test, Y_true_val, Y_true_test, Y_mask_val,  
      ↪Y_mask_test = train_test_split(split_test, X_test, Y_true_test, Y_mask_test,  
      ↪test_size=0.5)  
  
      train_set = np.asarray(set(split_train))  
      test_set = np.asarray(set(split_test))  
      print('# Train: {}, # Val: {}, # Test: {}'.format(len(split_train),  
      ↪len(split_test), len(split_val)))
```

```
# Train: 15000, # Val: 2500, # Test: 2500
```

```
[23]: print(X_train.shape)  
      print(X_val.shape)  
      print(X_test.shape)
```

```
(15000, 10, 128)
```

```
(2500, 10, 128)
```

```
(2500, 10, 128)
```

```
[29]: THRESHOLD = 0.3  
  
      # This dictionary will include the classifiers for each model  
      mymodels = dict()  
  
      # We'll iterate over all instrument classes, and fit a model for each one  
      # After training, we'll print a classification report for each instrument  
      for instrument in class_map:  
  
          # Map the instrument name to its column number  
          inst_num = class_map[instrument]
```

```

# Step 1: sub-sample the data

# First, we need to select down to the data for which we have annotations
# This is what the mask arrays are for
train_inst = Y_mask_train[:, inst_num]
val_inst = Y_mask_val[:, inst_num]
test_inst = Y_mask_test[:, inst_num]

# Here, we're using the Y_mask_train array to slice out only the training
→examples
# for which we have annotations for the given class
X_train_inst = X_train[train_inst]
X_val_inst = X_val[val_inst]

# Step 3: simplify the data by averaging over time

# Let's arrange the data for a sklearn Random Forest model
# Instead of having time-varying features, we'll summarize each track by
→its mean feature vector over time
X_train_inst_sklearn = np.mean(X_train_inst, axis=1)

# Again, we slice the labels to the annotated examples
# We threshold the label likelihoods at 0.5 to get binary labels
Y_true_train_inst = Y_true_train[train_inst, inst_num] >= THRESHOLD
Y_true_val_inst = Y_true_val[val_inst, inst_num] >= THRESHOLD

# Repeat the above slicing and dicing but for the test set
X_test_inst = X_test[test_inst]
X_test_inst_sklearn = np.mean(X_test_inst, axis=1)
Y_true_test_inst = Y_true_test[test_inst, inst_num] >= THRESHOLD

X_train_inst = X_train_inst.astype('float32')
X_val_inst = X_val_inst.astype('float32')
X_train_inst_sklearn = X_train_inst_sklearn.astype('float32')
X_train_inst_sklearn = lb.util.normalize(X_train_inst_sklearn)
# X_train_inst = S_dB
print(X_train_inst.shape)
shape = X_train_inst.shape
X_train_inst = X_train_inst.reshape(shape[0],1, shape[1], shape[2])
shape = X_val_inst.shape
X_val_inst = X_val_inst.reshape(shape[0],1, shape[1], shape[2])
shape = X_test_inst.shape
X_test_inst = X_test_inst.reshape(shape[0],1, shape[1], shape[2])
#X_train_inst = X_train_inst.reshape(1,1,431,128)
print(X_train_inst.shape)

```

```

print(Y_true_train_inst[0])
# Step 3.
# Initialize a new classifier
import keras, os
from keras.models import Sequential
from keras.layers import Dense, Conv2D, MaxPool2D, Flatten
from keras.preprocessing.image import ImageDataGenerator
import numpy as np
model = models.Sequential()

# model.add(layers.Conv2D(filters=8, kernel_size=(3,3), activation='relu',
→input_shape=(10,128,1,)))
model.
→add(Conv2D(input_shape=(1,10,128), data_format="channels_first", filters=64, kernel_size=(3,3)
→activation="relu"))
model.add(Conv2D(filters=32, kernel_size=(3,3), padding="same",
→activation="relu"))
model.add(MaxPool2D(pool_size=(3,3), strides=(2,2)))
model.add(Conv2D(filters=128, kernel_size=(3,3), padding="same",
→activation="relu"))
model.add(Conv2D(filters=128, kernel_size=(3,3), padding="same",
→activation="relu"))
model.add(MaxPool2D(pool_size=(2,2), strides=(2,2)))
model.add(Conv2D(filters=256, kernel_size=(3,3), padding="same",
→activation="relu"))
model.add(Conv2D(filters=256, kernel_size=(3,3), padding="same",
→activation="relu"))
model.add(Conv2D(filters=256, kernel_size=(3,3), padding="same",
→activation="relu"))
model.add(MaxPool2D(pool_size=(2,2), strides=(2,2)))
model.add(Conv2D(filters=512, kernel_size=(3,3), padding="same",
→activation="relu"))
model.add(Conv2D(filters=512, kernel_size=(3,3), padding="same",
→activation="relu"))
model.add(Conv2D(filters=512, kernel_size=(3,3), padding="same",
→activation="relu"))
model.add(layers.Flatten())
model.add(layers.Dense(units=4096, activation='relu'))
model.add(layers.Dense(units=4096, activation='relu'))
model.add(layers.Dense(units=1, activation='sigmoid'))

model.compile(loss='binary_crossentropy',
              optimizer=keras.optimizers.Adam(lr=0.0001),
              metrics = ['accuracy'])

# model.summary()

```

```

# Step 4.
history = model.fit(X_train_inst, Y_true_train_inst , epochs=10,
→batch_size=64, validation_data=(X_val_inst, Y_true_val_inst))

plt.figure(figsize=(9,4))

plt.subplot(1,2,1)
plt.plot(history.history['acc'])
plt.plot(history.history['val_acc'])
plt.title('Model accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train accuracy', 'Validation accuracy'], loc='upper left')

plt.subplot(1,2,2)
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train loss', 'Validation loss'], loc='upper left')
plt.show()

loss, acc = model.evaluate(X_test_inst, Y_true_test_inst)
print('Test loss: {}'.format(loss))
print('Test accuracy: {:.2%}'.format(acc))

# Step 5.
# Finally, we'll evaluate the model on both train and test
Y_pred_train = model.predict(X_train_inst)
Y_pred_test = model.predict(X_test_inst)
Y_pred_train_bool = Y_pred_train > THRESHOLD - 0.2 #THRESHOLD (should be
→lower than 0.5)
Y_pred_test_bool = Y_pred_test > THRESHOLD - 0.2 #THRESHOLD (should be
→lower than 0.5)
print(Y_pred_train[0])
print('-' * 52)
print(instrument)
print('\tTRAIN')
print(classification_report(Y_true_train_inst, Y_pred_train_bool))

print(Y_true_train_inst[3])
print(Y_pred_train[3])
print('\tTEST')
print(classification_report(Y_true_test_inst, Y_pred_test_bool))
sum = 0
# for i, prob in enumerate(Y_pred_train):

```

```

#     print (i)
#     print (prob)
#     sum += prob
#     print(sum)
# Store the classifier in our dictionary
mymodels[instrument] = model

```

(1566, 10, 128)

(1566, 1, 10, 128)

False

Train on 1566 samples, validate on 249 samples

Epoch 1/10

1566/1566 [=====] - 92s 58ms/step - loss: 0.5582 - acc: 0.7625 - val\_loss: 0.5340 - val\_acc: 0.7671

Epoch 2/10

1566/1566 [=====] - 95s 60ms/step - loss: 0.4559 - acc: 0.7842 - val\_loss: 0.4145 - val\_acc: 0.8233

Epoch 3/10

1566/1566 [=====] - 90s 58ms/step - loss: 0.3943 - acc: 0.8225 - val\_loss: 0.3986 - val\_acc: 0.8313

Epoch 4/10

1566/1566 [=====] - 89s 57ms/step - loss: 0.3811 - acc: 0.8276 - val\_loss: 0.4120 - val\_acc: 0.7831

Epoch 5/10

1566/1566 [=====] - 89s 57ms/step - loss: 0.3109 - acc: 0.8621 - val\_loss: 0.4202 - val\_acc: 0.8233

Epoch 6/10

1566/1566 [=====] - 88s 56ms/step - loss: 0.2811 - acc: 0.8838 - val\_loss: 0.4195 - val\_acc: 0.7992

Epoch 7/10

1566/1566 [=====] - 87s 56ms/step - loss: 0.2400 - acc: 0.9023 - val\_loss: 0.4760 - val\_acc: 0.7952

Epoch 8/10


1566/1566 [=====] - 89s 57ms/step - loss: 0.1566 - acc: 0.9374 - val\_loss: 0.6051 - val\_acc: 0.7912

Epoch 9/10

1566/1566 [=====] - 89s 57ms/step - loss: 0.1246 - acc: 0.9508 - val\_loss: 0.6610 - val\_acc: 0.8032

Epoch 10/10

1566/1566 [=====] - 88s 56ms/step - loss: 0.0765 - acc: 0.9764 - val\_loss: 0.5946 - val\_acc: 0.8112



256/256 [=====] - 2s 9ms/step

Test loss: 0.7445437833666801

Test accuracy: 75.00%

[0.03150272]

-----  
accordion

TRAIN				
	precision	recall	f1-score	support
False	1.00	0.98	0.99	1194
True	0.94	0.99	0.97	372
accuracy			0.98	1566
macro avg	0.97	0.99	0.98	1566
weighted avg	0.98	0.98	0.98	1566

False

[0.01145008]

TEST				
	precision	recall	f1-score	support
False	0.89	0.77	0.82	197
True	0.47	0.68	0.55	59
accuracy			0.75	256
macro avg	0.68	0.72	0.69	256
weighted avg	0.79	0.75	0.76	256


(1669, 10, 128)



```

(1669, 1, 10, 128)
True
Train on 1669 samples, validate on 279 samples
Epoch 1/10
1669/1669 [=====] - 100s 60ms/step - loss: 0.6433 -
acc: 0.6465 - val_loss: 0.6155 - val_acc: 0.6559
Epoch 2/10
1669/1669 [=====] - 96s 57ms/step - loss: 0.4962 - acc:
0.7358 - val_loss: 0.5193 - val_acc: 0.7491
Epoch 3/10
1669/1669 [=====] - 94s 57ms/step - loss: 0.4215 - acc:
0.7921 - val_loss: 0.4306 - val_acc: 0.8029
Epoch 4/10
1669/1669 [=====] - 95s 57ms/step - loss: 0.4049 - acc:
0.7999 - val_loss: 0.4494 - val_acc: 0.7993
Epoch 5/10
1669/1669 [=====] - 95s 57ms/step - loss: 0.4005 - acc:
0.8089 - val_loss: 0.4588 - val_acc: 0.7563
Epoch 6/10
1669/1669 [=====] - 94s 56ms/step - loss: 0.3557 - acc:
0.8340 - val_loss: 0.4518 - val_acc: 0.7706
Epoch 7/10
1669/1669 [=====] - 94s 56ms/step - loss: 0.3137 - acc:
0.8616 - val_loss: 0.4496 - val_acc: 0.7491
Epoch 8/10
1669/1669 [=====] - 98s 59ms/step - loss: 0.3010 - acc:
0.8634 - val_loss: 0.4336 - val_acc: 0.7814
Epoch 9/10
1669/1669 [=====] - 101s 61ms/step - loss: 0.2698 -
acc: 0.8814 - val_loss: 0.5598 - val_acc: 0.7670
Epoch 10/10
1669/1669 [=====] - 93s 56ms/step - loss: 0.2312 - acc:
0.9005 - val_loss: 0.5422 - val_acc: 0.7921

```



270/270 [=====] - 3s 10ms/step  
 Test loss: 0.5628210295129705  
 Test accuracy: 79.63%  
 [0.93295425]

-----  
 banjo

TRAIN				
	precision	recall	f1-score	support
False	0.97	0.85	0.91	1123
True	0.76	0.95	0.84	546
accuracy			0.89	1669
macro avg	0.87	0.90	0.88	1669
weighted avg	0.90	0.89	0.89	1669

False

[0.00056961]

TEST				
	precision	recall	f1-score	support
False	0.87	0.73	0.80	180
True	0.59	0.78	0.67	90
accuracy			0.75	270
macro avg	0.73	0.76	0.73	270
weighted avg	0.78	0.75	0.75	270

(1401, 10, 128)

```

(1401, 1, 10, 128)
False
Train on 1401 samples, validate on 258 samples
Epoch 1/10
1401/1401 [=====] - 87s 62ms/step - loss: 0.6153 - acc:
0.7059 - val_loss: 0.5531 - val_acc: 0.7364
Epoch 2/10
1401/1401 [=====] - 78s 56ms/step - loss: 0.5357 - acc:
0.7281 - val_loss: 0.4700 - val_acc: 0.7674
Epoch 3/10
1401/1401 [=====] - 83s 59ms/step - loss: 0.4440 - acc:
0.7844 - val_loss: 0.4689 - val_acc: 0.7791
Epoch 4/10
960/1401 [=====>...] - ETA: 25s - loss: 0.3998 - acc:

```

```

KeyboardInterrupt                                Traceback (most recent call
last)

<ipython-input-29-b6a7d116500e> in <module>
    91     # model.summary()
    92     # Step 4.
--> 93     history = model.fit(X_train_inst,Y_true_train_inst , epochs=10,
batch_size=64, validation_data=(X_val_inst,Y_true_val_inst))
    94
    95     plt.figure(figsize=(9,4))

~\Documents\Conda\lib\site-packages\keras\engine\training.py in
fit(self, x, y, batch_size, epochs, verbose, callbacks, validation_split,
validation_data, shuffle, class_weight, sample_weight, initial_epoch,
steps_per_epoch, validation_steps, **kwargs)
    1037                                     initial_epoch=initial_epoch,
    1038
-> steps_per_epoch=steps_per_epoch,
-> 1039
validation_steps=validation_steps)
    1040
    1041     def evaluate(self, x=None, y=None,

~\Documents\Conda\lib\site-packages\keras\engine\training_arrays.py in
fit_loop(model, f, ins, out_labels, batch_size, epochs, verbose, callbacks,
val_f, val_ins, shuffle, callback_metrics, initial_epoch, steps_per_epoch,
validation_steps)

```

```

197             ins_batch[i] = ins_batch[i].toarray()
198
--> 199         outs = f(ins_batch)
200         outs = to_list(outs)
201         for l, o in zip(out_labels, outs):

~\Documents\Conda\lib\site-packages\keras\backend\tensorflow_backend.py
↵in __call__(self, inputs)
2713             return self._legacy_call(inputs)
2714
-> 2715         return self._call(inputs)
2716     else:
2717         if py_any(is_tensor(x) for x in inputs):

~\Documents\Conda\lib\site-packages\keras\backend\tensorflow_backend.py
↵in _call(self, inputs)
2673         fetched = self._callable_fn(*array_vals,
↵run_metadata=self.run_metadata)
2674     else:
-> 2675         fetched = self._callable_fn(*array_vals)
2676         return fetched[:len(self.outputs)]
2677

~\Documents\Conda\lib\site-packages\tensorflow\python\client\session.py
↵in __call__(self, *args, **kwargs)
1437         ret = tf_session.TF_SessionRunCallable(
1438             self._session._session, self._handle, args, status,
-> 1439             run_metadata_ptr)
1440         if run_metadata:
1441             proto_data = tf_session.TF_GetBuffer(run_metadata_ptr)

```

KeyboardInterrupt:

```

[: print(X_train_inst_sklearn)
print(Y_pred_train)

```

```

[:

```